

```
In [3]: ## Madhavi Ghanta  
## DSC 680 Project 1  
## Credit Card Fraud
```

```
In [4]: ### Load necessary Libraries
```

```
In [ ]: import pandas as pd  
import numpy as np  
import plotly.express as px  
import seaborn as sns  
import matplotlib.pyplot as plt  
from matplotlib.ticker import NullFormatter  
import opendatasets as od  
  
from sklearn.model_selection import train_test_split,cross_val_score  
from sklearn.preprocessing import StandardScaler, LabelEncoder  
from sklearn.feature_selection import chi2, SelectKBest  
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score,confusion_matrix,  
from imblearn.over_sampling import SMOTE  
  
from sklearn.linear_model import LogisticRegression  
from sklearn.ensemble import RandomForestClassifier  
from xgboost import XGBClassifier  
  
import tensorflow as tf  
from tensorflow import keras  
from tensorflow.keras import layers
```

Dataset is already split into test and train sets. We will combine them and redo the train-test split

```
In [25]: fraud_train_df = pd.read_csv("C:/Users/mghan/Documents/MSDS/DSC680/Project1/fraudTrain  
fraud_train_df.head(5)
```

Out[25]:

| | Unnamed: 0 | trans_date_trans_time | cc_num | merchant | category | amt | first |
|----------|-----------------------|------------------------------|------------------|--|-----------------|------------|--------------|
| 0 | 0 | 2019-01-01 00:00:18 | 2703186189652095 | fraud_Rippin, Kub and Mann | misc_net | 4.97 | Jennifer |
| 1 | 1 | 2019-01-01 00:00:44 | 630423337322 | fraud_Heller, Gutmann and Zieme | grocery_pos | 107.23 | Stephanie |
| 2 | 2 | 2019-01-01 00:00:51 | 38859492057661 | fraud_Lind- Buckridge | entertainment | 220.11 | Edward |
| 3 | 3 | 2019-01-01 00:01:16 | 3534093764340240 | fraud_Kutch, Hermiston and Farrell | gas_transport | 45.00 | Jeremy |
| 4 | 4 | 2019-01-01 00:03:06 | 375534208663984 | fraud_Keeling- Crist | misc_pos | 41.96 | Tyler |

5 rows × 23 columns

In [24]:

```
fraud_test_df = pd.read_csv("C:/Users/mghan/Documents/MSDS/DSC680/Project1/fraudTest.csv")
fraud_test_df.head(5)
```

Out[24]:

| | Unnamed: 0 | trans_date_trans_time | cc_num | merchant | category | amt | first |
|---|---------------|-----------------------|------------------|--------------------------------------|----------------|-------|--------|
| 0 | 0 | 2020-06-21 12:14:25 | 2291163933867244 | fraud_Kirlin and Sons | personal_care | 2.86 | Jeff |
| 1 | 1 | 2020-06-21 12:14:33 | 3573030041201292 | fraud_Sporer-Keebler | personal_care | 29.84 | Joanne |
| 2 | 2 | 2020-06-21 12:14:53 | 3598215285024754 | fraud_Swaniawski, Nitzsche and Welch | health_fitness | 41.28 | Ashley |
| 3 | 3 | 2020-06-21 12:15:15 | 3591919803438423 | fraud_Haley Group | misc_pos | 60.05 | Brian |
| 4 | 4 | 2020-06-21 12:15:17 | 3526826139003047 | fraud_Johnston-Casper | travel | 3.19 | Nathan |

5 rows × 23 columns



In [26]:

```
#Combine the 2 datasets
fraud_df = pd.concat([fraud_train_df, fraud_test_df], axis=0)
```

In [27]:

```
fraud_train_df.shape, fraud_test_df.shape, fraud_df.shape
```

Out[27]:

```
((1296675, 23), (555719, 23), (1852394, 23))
```

In [28]:

```
fraud_df.head(5)
```

Out[28]:

| | 0 | trans_date_trans_time | cc_num | merchant | category | amt | first |
|---|---|-----------------------|------------------|--|---------------|--------|-----------|
| 0 | 0 | 2019-01-01 00:00:18 | 2703186189652095 | fraud_Rippin, Kub and Mann | misc_net | 4.97 | Jennifer |
| 1 | 1 | 2019-01-01 00:00:44 | 630423337322 | fraud_Heller, Gutmann and Zieme | grocery_pos | 107.23 | Stephanie |
| 2 | 2 | 2019-01-01 00:00:51 | 38859492057661 | fraud_Lind- Buckridge | entertainment | 220.11 | Edward |
| 3 | 3 | 2019-01-01 00:01:16 | 3534093764340240 | fraud_Kutch, Hermiston and Farrell | gas_transport | 45.00 | Jeremy |
| 4 | 4 | 2019-01-01 00:03:06 | 375534208663984 | fraud_Keeling- Crist | misc_pos | 41.96 | Tyler |

5 rows × 23 columns



Data Processing

Null rows check

In [29]:

```
fraud_df.isnull().sum()
```

```
Out[29]: Unnamed: 0          0
         trans_date_trans_time 0
         cc_num                0
         merchant              0
         category              0
         amt                   0
         first                 0
         last                  0
         gender                0
         street                0
         city                  0
         state                 0
         zip                   0
         lat                   0
         long                  0
         city_pop              0
         job                   0
         dob                   0
         trans_num              0
         unix_time              0
         merch_lat              0
         merch_long             0
         is_fraud               0
         dtype: int64
```

Check for duplicates

```
In [30]: print('Dataframe before dropping duplicates :', fraud_df.shape)
fraud_df = fraud_df.drop_duplicates() # 1,389 rows dropped
print('Dataframe after dropping duplicates :', fraud_df.shape)
```

```
Dataframe before dropping duplicates : (1852394, 23)
Dataframe after dropping duplicates : (1852394, 23)
```

```
In [31]: fraud_df.duplicated().sum()
```

```
Out[31]: 0
```

Convert column data type for DateTime

```
In [32]: fraud_df['trans_date_time'] = pd.to_datetime(fraud_df['trans_date_trans_time'])
fraud_df.head(5)
```

Out[32]:

| | Unnamed: 0 | trans_date_trans_time | cc_num | merchant | category | amt | first |
|----------|-----------------------|------------------------------|------------------|--|-----------------|------------|--------------|
| 0 | 0 | 2019-01-01 00:00:18 | 2703186189652095 | fraud_Rippin, Kub and Mann | misc_net | 4.97 | Jennifer |
| 1 | 1 | 2019-01-01 00:00:44 | 630423337322 | fraud_Heller, Gutmann and Zieme | grocery_pos | 107.23 | Stephanie |
| 2 | 2 | 2019-01-01 00:00:51 | 38859492057661 | fraud_Lind- Buckridge | entertainment | 220.11 | Edward |
| 3 | 3 | 2019-01-01 00:01:16 | 3534093764340240 | fraud_Kutch, Hermiston and Farrell | gas_transport | 45.00 | Jeremy |
| 4 | 4 | 2019-01-01 00:03:06 | 375534208663984 | fraud_Keeling- Crist | misc_pos | 41.96 | Tyler |

5 rows × 24 columns



Add columns

```
In [33]: fraud_df['month'] = pd.DatetimeIndex(fraud_df['trans_date_time']).month
```

Drop Columns

```
In [34]: fraud_df.columns
```

```
Out[34]: Index(['Unnamed: 0', 'trans_date_trans_time', 'cc_num', 'merchant', 'category',  
       'amt', 'first', 'last', 'gender', 'street', 'city', 'state', 'zip',  
       'lat', 'long', 'city_pop', 'job', 'dob', 'trans_num', 'unix_time',  
       'merch_lat', 'merch_long', 'is_fraud', 'trans_date_time', 'month'],  
      dtype='object')
```

```
In [35]: fraud_df = fraud_df.drop(['Unnamed: 0', 'trans_date_trans_time'], axis=1)
```

```
In [36]: fraud_df.columns
```

```
Out[36]: Index(['cc_num', 'merchant', 'category', 'amt', 'first', 'last', 'gender',  
       'street', 'city', 'state', 'zip', 'lat', 'long', 'city_pop', 'job',  
       'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long', 'is_fraud',  
       'trans_date_time', 'month'],  
      dtype='object')
```

Convert data types to reduce memory usage

```
In [37]: fraud_df.dtypes
```

```
Out[37]: cc_num           int64
merchant          object
category          object
amt             float64
first            object
last             object
gender            object
street            object
city              object
state             object
zip                int64
lat               float64
long              float64
city_pop          int64
job               object
dob               object
trans_num         object
unix_time         int64
merch_lat         float64
merch_long        float64
is_fraud          int64
trans_date_time   datetime64[ns]
month             int64
dtype: object
```

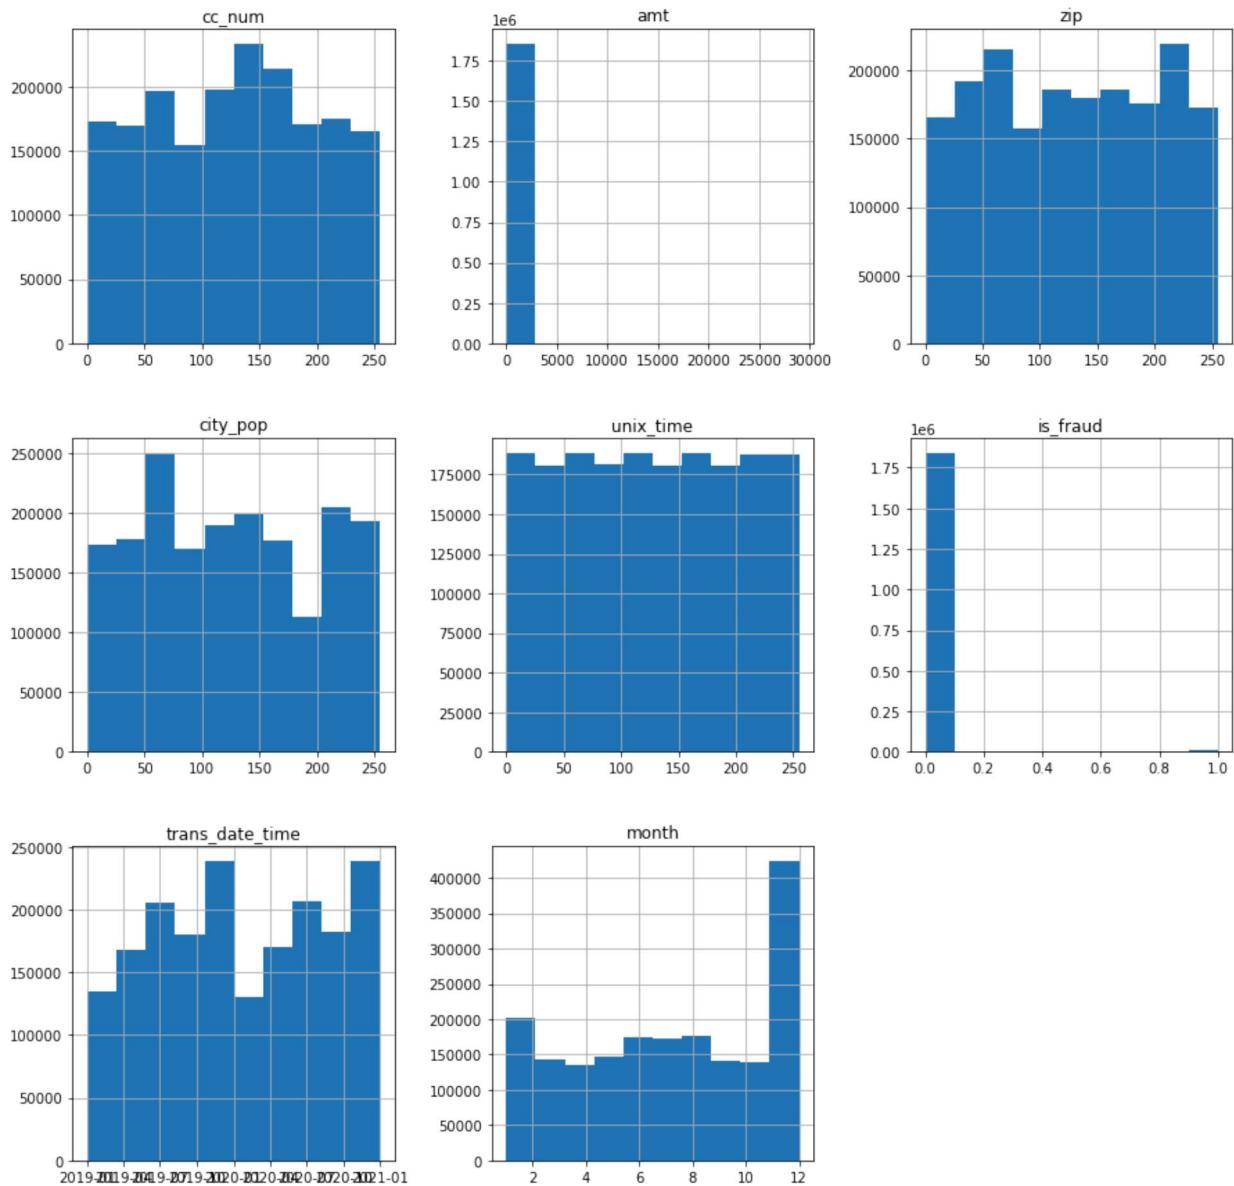
```
In [38]: #Converting data types to avoid memory issues while executing the model fit.
cols=['amt']
fraud_df[cols] = fraud_df[cols].astype('float16') #Converting float64 to float16

int_cols = ['cc_num','zip','city_pop','unix_time','is_fraud','month']
fraud_df[int_cols] = fraud_df[int_cols].astype(np.uint8) #Converting int64 to uint8

obj_cols = ['merchant','category','first','last','gender','street','city','state','job'
fraud_df[obj_cols] = fraud_df[obj_cols].astype('category') # #Converting object to cat
```

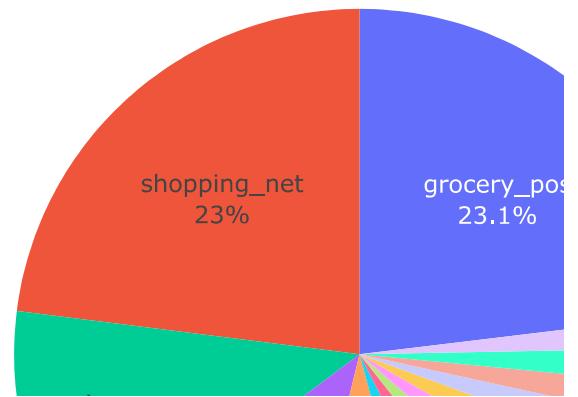
Data Visualization

```
In [39]: fraud_df.hist(figsize = (15,15))
plt.show()
```



```
In [40]: fig = px.pie(fraud_df[fraud_df.is_fraud==1], values='is_fraud', names='category', title="Percentage of Fraud by Category")
fig.update_traces(textposition='inside', textinfo='percent+label')
fig.update_layout(title = "Percentage of Fraud by Category")
fig.show("notebook")
```

Percentage of Fraud by Category



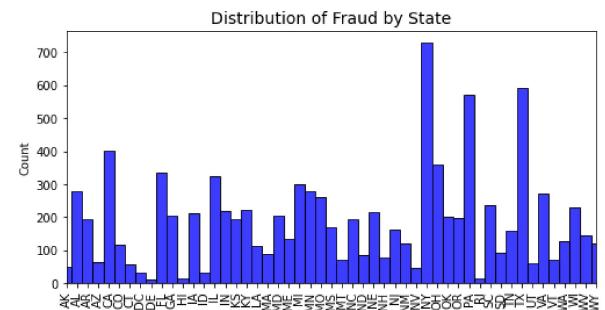
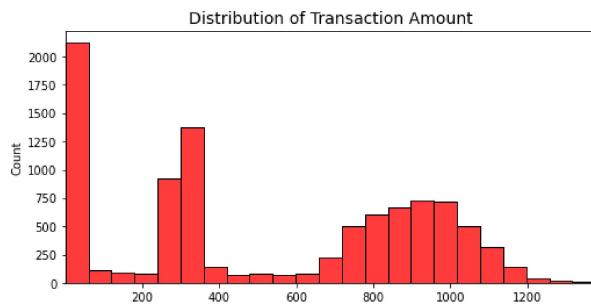
```
In [41]: fig, ax = plt.subplots(1, 2, figsize=(18,4))

amount_val = fraud_df[ fraud_df.is_fraud == 1].amt.values.astype(int)
time_val = fraud_df[ fraud_df.is_fraud == 1].state.values

sns.histplot(amount_val, ax=ax[0], color='r')
ax[0].set_title('Distribution of Transaction Amount', fontsize=14)
ax[0].set_xlim([min(amount_val), max(amount_val)])

sns.histplot(time_val, ax=ax[1], color='b')
ax[1].set_title('Distribution of Fraud by State', fontsize=14)
ax[1].set_xlim([min(time_val), max(time_val)])
ax[1].tick_params(axis='x', rotation=90)

plt.show()
```



We can see that the data is not balanced. The number of fraudulent transactions in the dataset are very low in comparison to the legitimate transactions. Building models with this data could give inaccurate results.

Model Building

SMOTE helps to balance the class distribution by generating synthetic samples of the minority class.

Data leakage happens when information from the validation or test set unintentionally leaks into the training set, leading to overly optimistic performance estimates. When using SMOTE, this risk exists because synthetic samples are generated based on the original data. By incorporating SMOTE within the cross-validation process, you mitigate the risk of data leakage, as the synthetic samples are only used within each fold of the cross-validation.

```
In [45]: fraud_df.replace(np.nan,0)
```

Out[45]:

| | cc_num | merchant | category | amt | first | last | gender | street |
|--------|--------|------------------------------------|----------------|------------|-----------|---------|--------|------------------------------|
| 0 | 127 | fraud_Rippin, Kub and Mann | misc_net | 4.968750 | Jennifer | Banks | F | 561 Perry Cove |
| 1 | 106 | fraud_Heller, Gutmann and Zieme | grocery_pos | 107.250000 | Stephanie | Gill | F | 43039 Riley Greens Suite 393 |
| 2 | 61 | fraud_Lind-Buckridge | entertainment | 220.125000 | Edward | Sanchez | M | 594 White Dale Suite 530 |
| 3 | 16 | fraud_Kutch, Hermiston and Farrell | gas_transport | 45.000000 | Jeremy | White | M | 9443 Cynthia Court Apt. 038 |
| 4 | 176 | fraud_Keeling-Crist | misc_pos | 41.968750 | Tyler | Garcia | M | 408 Bradley Rest |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 555714 | 169 | fraud_Reilly and Sons | health_fitness | 43.781250 | Michael | Olson | M | 558 Michael Estates |
| 555715 | 40 | fraud_Hoppe-Parisian | kids_pets | 111.812500 | Jose | Vasquez | M | 572 Davis Mountains |
| 555716 | 230 | fraud_Rau-Robel | kids_pets | 86.875000 | Ann | Lawson | F | 144 Evans Islands Apt. 683 |
| 555717 | 150 | fraud_Breitenberg LLC | travel | 7.988281 | Eric | Preston | M | 7020 Doyle Stream Apt. 951 |
| 555718 | 187 | fraud_Dare-Marvin | entertainment | 38.125000 | Samuel | Frey | M | 830 Myers Plaza Apt. 384 |

1852394 rows × 23 columns



In [46]:

```
# Split the dataset into train and test sets
#X = fraud_df.drop(['is_fraud', 'lat', 'Long', 'merchLat', 'merchLong'], axis=1)
X = fraud_df.drop(['is_fraud'], axis=1)
Y = fraud_df['is_fraud']
X.shape, Y.shape
```

Out[46]:

```
((1852394, 22), (1852394,))
```

In [47]:

```
# Encode categorical variables (e.g., 'gender', 'category', 'state', etc.)
categorical_columns = [ 'merchant', 'category', 'first', 'last', 'gender', 'street', 'city']
for col in categorical_columns:
```

```

le = LabelEncoder()
X[col] = le.fit_transform(X[col])

```

In [48]: # Standardize numerical features
scaler = StandardScaler()
X = scaler.fit_transform(X)

In [49]: # Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size=0.2, random_state=42)

In [50]: X_train.shape, Y_train.shape
Out[50]: ((1481915, 22), (1481915,))

In [51]: # Apply SMOTE to balance the dataset
smote = SMOTE(sampling_strategy='auto', random_state=42)
x_train_resampled, y_train_resampled = smote.fit_resample(X_train, Y_train)

In [52]: print('x_train Data Shape : ', X_train.shape)
print('y_train Labels Shape : ', Y_train.shape)
print('x_train_resampled Data Shape : ', x_train_resampled.shape)
print('y_train_resampled Labels Shape : ', y_train_resampled.shape)
print('x_test Data Shape : ', X_test.shape)
print('y_test Labels Shape : ', Y_test.shape)

```

x_train Data Shape : (1481915, 22)
y_train Labels Shape : (1481915,)
x_train_resampled Data Shape : (2948434, 22)
y_train_resampled Labels Shape : (2948434,)
x_test Data Shape : (370479, 22)
y_test Labels Shape : (370479,)

```

In [53]: #Adding this step to clear memory, to avoid memory issues during execution
import gc

gc.collect()

Out[53]: 35676

Models

RandomForestClassifier

In [54]: # Use the RandomForestClassifier to fit balanced data
rfc = RandomForestClassifier()
rfc_model = rfc.fit(x_train_resampled,y_train_resampled)

#Predict y data with classifier:
y_pred_rfc = rfc_model.predict(X_test)

Evaluate the model
print(classification_report(Y_test, y_pred_rfc))
print(confusion_matrix(Y_test, y_pred_rfc))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_rfc)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_rfc)}')

```

precision    recall   f1-score   support
          0      1.00      1.00      1.00     368526
          1      0.78      0.72      0.75      1953

accuracy                           1.00     370479
macro avg       0.89      0.86      0.87     370479
weighted avg    1.00      1.00      1.00     370479

[[368136    390]
 [  548   1405]]
ROC-AUC score : 0.8591738860069383
Accuracy score : 0.9974681425937773

```

```

In [55]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_rfc, labels=[1,0])

print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

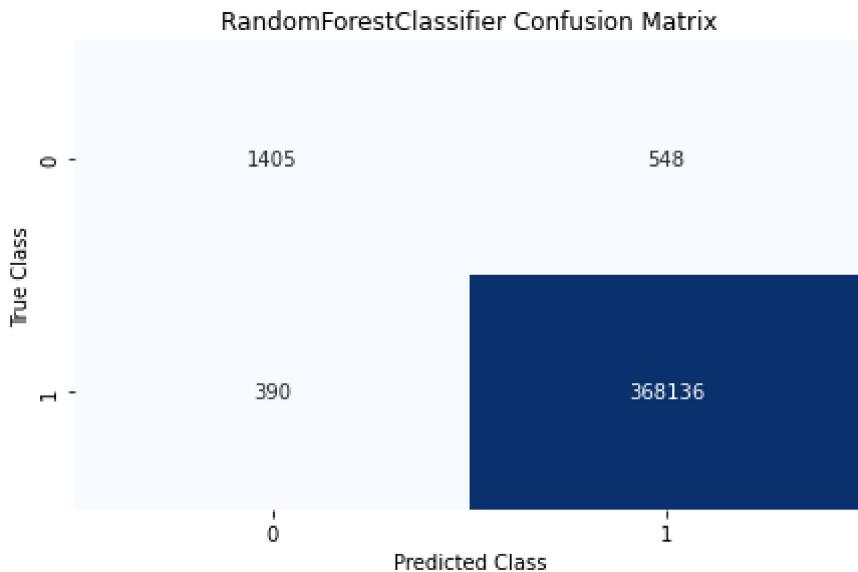
# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues",fmt='.{0f}')
plt.title("RandomForestClassifier Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()

```

```

[[ 1405    548]
 [ 390 368136]]

```



Logistic Regression

```

In [56]: # Train a logistic regression model
logistic_model = LogisticRegression(solver='liblinear', random_state=42)
logistic_model.fit(x_train_resampled,y_train_resampled)

# Make predictions on the test set
y_pred_lr = logistic_model.predict(X_test)

# Evaluate the model

```

```

print(classification_report(Y_test, y_pred_lr))
print(confusion_matrix(Y_test, y_pred_lr))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_lr)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_lr)}')

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.94 | 0.97 | 368526 |
| 1 | 0.06 | 0.77 | 0.12 | 1953 |
| accuracy | | | 0.94 | 370479 |
| macro avg | 0.53 | 0.85 | 0.54 | 370479 |
| weighted avg | 0.99 | 0.94 | 0.96 | 370479 |

```

[[346742  21784]
 [ 456   1497]]
ROC-AUC score : 0.8537009475361442
Accuracy score : 0.939969606914292

```

```

In [57]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_lr, labels=[1,0])

print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues", fmt=' .0f ')
plt.title("LogisticRegression Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()

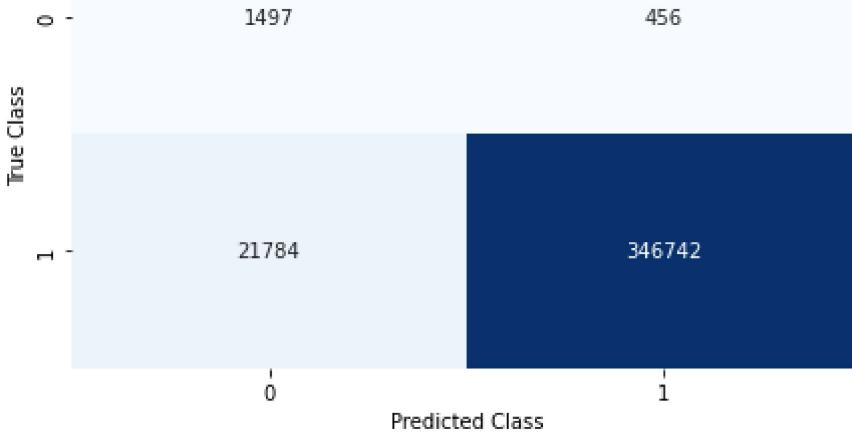
```

```

[[ 1497    456]
 [ 21784  346742]]

```

LogisticRegression Confusion Matrix



Using `class_weight='balanced'` to check if the imbalance get's any better.

```

In [58]: # Train a logistic regression model
logit = LogisticRegression(solver='liblinear', class_weight='balanced')
model_logit = logit.fit(x_train_resampled, y_train_resampled)

```

```

# Make predictions on the test set
y_pred_logit = model_logit.predict(X_test)

# Evaluate the model
print(classification_report(Y_test, y_pred_logit))
print(confusion_matrix(Y_test, y_pred_logit))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_logit)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_logit)}')

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.94 | 0.97 | 368526 |
| 1 | 0.06 | 0.77 | 0.12 | 1953 |
| accuracy | | | 0.94 | 370479 |
| macro avg | 0.53 | 0.85 | 0.54 | 370479 |
| weighted avg | 0.99 | 0.94 | 0.96 | 370479 |

```

[[346742  21784]
 [ 456   1497]]
ROC-AUC score : 0.8537009475361442
Accuracy score : 0.939969606914292

```

In [59]:

```

#Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_logit, labels=[1,0])

print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues", fmt='.{0f}')
plt.title("LogisticRegression Confusion Matrix with class_weight=balanced"), plt.tight
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()

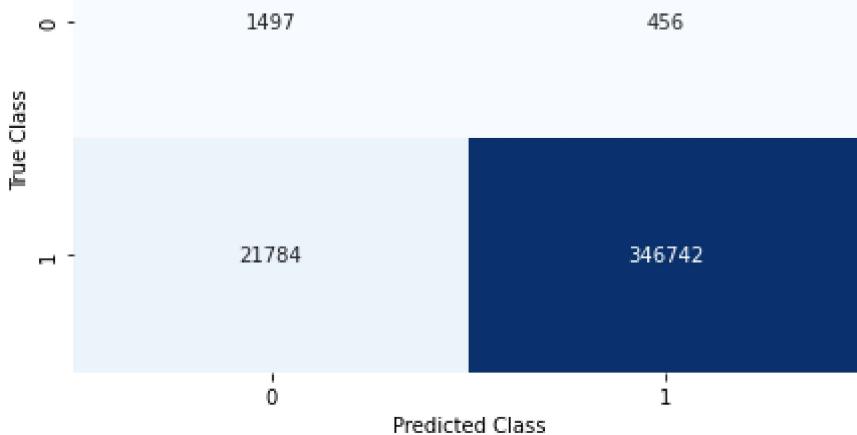
```

```

[[ 1497    456]
 [ 21784  346742]]

```

LogisticRegression Confusion Matrix with class_weight=balanced



Adding the class_weight = balanced has no impact on the model outcome.

Gradient Boosting

```
In [60]: xgb_model = XGBClassifier(max_depth = 4)
xgb_model.fit(x_train_resampled,y_train_resampled)
xgb_predicted = xgb_model.predict(X_test)
```

```
In [61]: # Train an XGBoost classifier
xgb_model = XGBClassifier(random_state=42)
xgb_model.fit(x_train_resampled,y_train_resampled)

# Make predictions on the test set
y_pred_gb = xgb_model.predict(X_test)

# Evaluate the XGBoost Classifier
print(classification_report(Y_test, y_pred_gb))
print(confusion_matrix(Y_test, y_pred_gb))
print(f'ROC-AUC score : {roc_auc_score(Y_test,y_pred_gb)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_gb)}')
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 1.00 | 1.00 | 368526 |
| 1 | 0.50 | 0.89 | 0.64 | 1953 |
| accuracy | | | 0.99 | 370479 |
| macro avg | 0.75 | 0.94 | 0.82 | 370479 |
| weighted avg | 1.00 | 0.99 | 1.00 | 370479 |

[[366819 1707]
 [217 1736]]
ROC-AUC score : 0.9421284613116396
Accuracy score : 0.9948067231880889

```
In [62]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_gb, labels=[1,0])

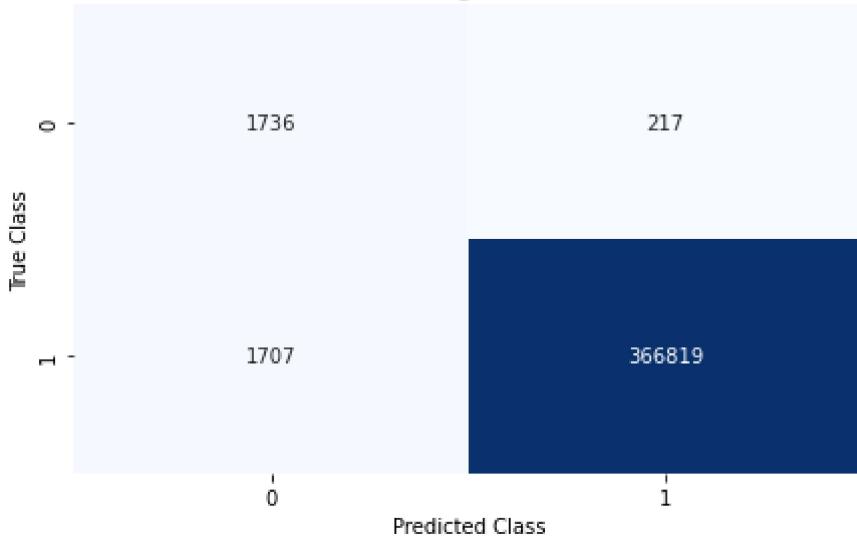
print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues", fmt='.\0f')
plt.title("Gradient Boosting Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()

[[ 1736   217]
 [ 1707 366819]]
```

Gradient Boosting Confusion Matrix



Neural Network Model

```
In [63]: # Build a neural network model
model = keras.Sequential([
    layers.Input(shape=(x_train_resampled.shape[1],)),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(32, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(1, activation='sigmoid')
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train the neural network
model.fit(x_train_resampled, y_train_resampled, epochs=10, batch_size=128, validation_
```

```
Epoch 1/10
18428/18428 [=====] - 43s 2ms/step - loss: 0.3090 - accuracy: 0.8902 - val_loss: 0.4048 - val_accuracy: 0.7722
Epoch 2/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2668 - accuracy: 0.9005 - val_loss: 0.3464 - val_accuracy: 0.7950
Epoch 3/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2553 - accuracy: 0.9026 - val_loss: 0.3220 - val_accuracy: 0.8016
Epoch 4/10
18428/18428 [=====] - 40s 2ms/step - loss: 0.2494 - accuracy: 0.9034 - val_loss: 0.3422 - val_accuracy: 0.8019
Epoch 5/10
18428/18428 [=====] - 42s 2ms/step - loss: 0.2471 - accuracy: 0.9037 - val_loss: 0.3302 - val_accuracy: 0.8026
Epoch 6/10
18428/18428 [=====] - 42s 2ms/step - loss: 0.2455 - accuracy: 0.9041 - val_loss: 0.3414 - val_accuracy: 0.8013
Epoch 7/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2444 - accuracy: 0.9041 - val_loss: 0.3097 - val_accuracy: 0.8033
Epoch 8/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2430 - accuracy: 0.9044 - val_loss: 0.3268 - val_accuracy: 0.8031
Epoch 9/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2421 - accuracy: 0.9044 - val_loss: 0.3172 - val_accuracy: 0.8030
Epoch 10/10
18428/18428 [=====] - 41s 2ms/step - loss: 0.2417 - accuracy: 0.9043 - val_loss: 0.3178 - val_accuracy: 0.8027
<keras.src.callbacks.History at 0x19a0988d460>
```

Out[63]:

```
In [64]: # Make predictions on the test set
y_pred_nn = model.predict(X_test)
y_pred_binary = (y_pred_nn > 0.5).astype(int)

# Evaluate the XGBoost Classifier
print(classification_report(Y_test, y_pred_binary))
print(confusion_matrix(Y_test, y_pred_binary))
print(f'ROC-AUC score : {roc_auc_score(Y_test,y_pred_binary)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_binary)}')

11578/11578 [=====] - 11s 975us/step
      precision    recall   f1-score   support
          0         1.00     0.99     0.99    368526
          1         0.29     0.85     0.43     1953

      accuracy                           0.99    370479
     macro avg       0.65     0.92     0.71    370479
  weighted avg       1.00     0.99     0.99    370479

[[364457  4069]
 [ 284 1669]]
ROC-AUC score : 0.9217707049546902
Accuracy score : 0.9882503461734673
```

```
In [65]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_binary, labels=[1,0])
```

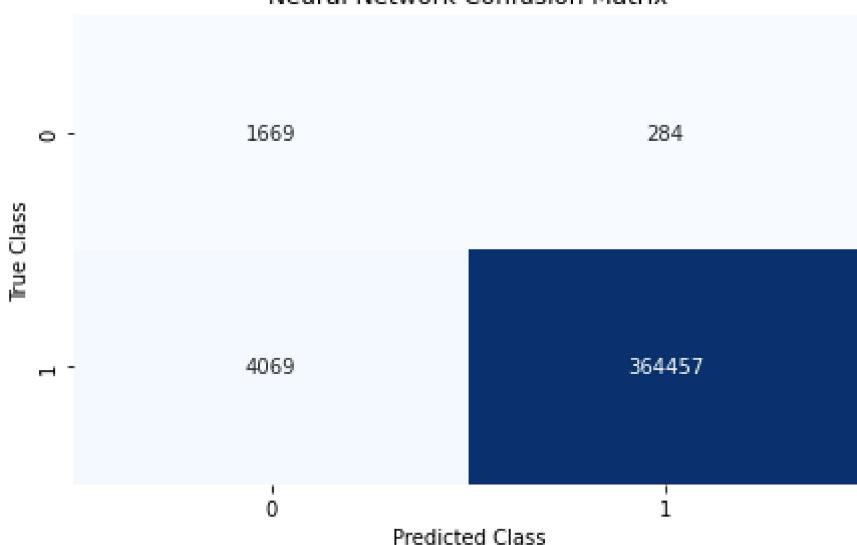
```
print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues", fmt=".0f")
plt.title("Neural Network Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()
```

```
[[ 1669   284]
 [ 4069 364457]]
```

Neural Network Confusion Matrix



In []: