

Seattle Weather

May 31, 2024

1 Has rain gotten more frequent in Seattle?

1.1 Madhavi Ghanta

<https://www.kaggle.com/rtatman/did-it-rain-in-seattle-19482017>

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from scipy.stats import binom
%matplotlib inline
```

```
[3]: rain = pd.read_csv("C:/Users/mghan/Documents/MSDS/Portfolio/
↪seattleWeather_1948-2017.csv")
```

```
[4]: rain.describe()
rain.head()
```

```
[4]:
```

	DATE	PRCP	TMAX	TMIN	RAIN
0	1/1/1948	0.47	51	42	True
1	1/2/1948	0.59	45	36	True
2	1/3/1948	0.42	45	35	True
3	1/4/1948	0.31	45	34	True
4	1/5/1948	0.17	45	32	True

```
[8]: #Bad Values within the data set
rain[pd.isnull(rain).any(axis=1)]
```

```
[8]:
```

	DATE	PRCP	TMAX	TMIN	RAIN
18415	6/2/1998	NaN	72	52	NaN
18416	6/3/1998	NaN	66	51	NaN
21067	9/5/2005	NaN	70	52	NaN

```
[9]: #Cleaning data set
rain.dropna(axis=1, how='all', inplace=True)
```

2 Hypothesis:

Rain has been steadily increasing in Seattle. Null Hypothesis: Rain has not increased or changed over time

2.1 Describe what the 5 variables mean in the dataset (Chapter 1).

DATE: the date of the observation

PRCP: the amount of precipitation, in inches

TMAX: the maximum temperature for that day, in degrees Fahrenheit

TMIN: the minimum temperature for that day, in degrees Fahrenheit

RAIN: TRUE if rain was observed on that day, FALSE if it was not


Include a histogram of each of the 5 variables – in your summary and analysis, identify any outliers and explain the reasoning for them being outliers and how you believe they should be handled (Chapter 2).

```
[10]: #Skipping date on histogram  
sns.distplot(rain['PRCP'].dropna())
```

```
C:\Users\mghan\anaconda3\lib\site-packages\seaborn\distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).
```

```
warnings.warn(msg, FutureWarning)
```

```
[10]: <AxesSubplot:xlabel='PRCP', ylabel='Density'>
```




```
[11]: sns.distplot(rain['TMAX'].dropna())
```

```
C:\Users\mgchan\anaconda3\lib\site-packages\seaborn\distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in a  
future version. Please adapt your code to use either `displot` (a figure-level  
function with similar flexibility) or `histplot` (an axes-level function for  
histograms).
```

```
warnings.warn(msg, FutureWarning)
```


```
[11]: <AxesSubplot:xlabel='TMAX', ylabel='Density'>
```



```
[12]: sns.distplot(rain['TMIN'].dropna())
```


```
C:\Users\mghan\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a
future version. Please adapt your code to use either `displot` (a figure-level
function with similar flexibility) or `histplot` (an axes-level function for
histograms).
  warnings.warn(msg, FutureWarning)
```

```
[12]: <AxesSubplot:xlabel='TMIN', ylabel='Density'>
```



```
[13]: rain['RAIN'].value_counts().plot(kind='bar', label='Rain')
```

```
[13]: <AxesSubplot:>
```



Include the other descriptive characteristics about the variables: * Mean * Mode * Spread * Tails

```
[14]: rain['TMIN'].mean(), rain['TMIN'].median(), rain['TMIN'].mode(), rain['TMIN'].  
      ↪std()
```

```
[14]: (44.51422644906266,  
      45.0,  
      0    42  
      Name: TMIN, dtype: int64,  
      8.892835742411922)
```

```
[15]: rain['DATE'].max(), rain['DATE'].min()
```

```
[15]: ('9/9/2017', '1/1/1948')
```

```
[16]: rain['DATE'] = pd.to_datetime(rain['DATE'])
```

```
[17]: rain['YEAR'] = rain['DATE'].dt.year
```

```
[18]: rain.head()
```


```
[18]:
```

	DATE	PRCP	TMAX	TMIN	RAIN	YEAR
0	1948-01-01	0.47	51	42	True	1948
1	1948-01-02	0.59	45	36	True	1948
2	1948-01-03	0.42	45	35	True	1948
3	1948-01-04	0.31	45	34	True	1948
4	1948-01-05	0.17	45	32	True	1948

```
[19]: first_half = rain[rain["DATE"]>='1982-12-22 00:00:00']  
      second_half = rain[rain["DATE"]<='1982-12-23 00:00:00']
```


```
[22]: plt.bar(list(pmf_first.values()),list(pmf_first.keys()))
```

```
[22]: <BarContainer object of 183 artists>
```



```
[25]: plt.bar(list(pmf_second.values()), list(pmf_second.keys()))
```

```
[25]: <BarContainer object of 175 artists>
```



[113]: pmf

```
[113]: {0.0: 7246,  
        0.01: 462,  
        0.02: 353,  
        0.03: 258,  
        0.04: 198,  
        0.05: 212,  
        0.06: 209,  
        0.07: 170,  
        0.08: 174,  
        0.09: 130,  
        0.1: 129,  
        0.11: 142,  
        0.12: 146,  
        0.13: 130,  
        0.14: 112,  
        0.15: 108,  
        0.16: 97,  
        0.17: 96,  
        0.18: 70,  
        0.19: 63,  
        0.2: 101,  
        0.21: 73,  
        0.22: 62,  
        0.23: 77,  
        0.24: 70,  
        0.25: 62,  
        0.26: 59,  
        0.27: 51,  
        0.28: 68,  
        0.29: 64,  
        0.3: 67,  
        0.31: 54,  
        0.32: 38,  
        0.33: 51,  
        0.34: 47,  
        0.35: 53,  
        0.36: 49,  
        0.37: 41,  
        0.38: 44,  
        0.39: 42,  
        0.4: 29,  
        0.41: 28,  
        0.42: 32,  
        0.43: 26,  
        0.44: 33,
```


0.45: 46,
0.46: 27,
0.47: 30,
0.48: 22,
0.49: 31,
0.5: 26,
0.51: 24,
0.52: 17,
0.53: 25,
0.54: 24,
0.55: 19,
0.56: 35,
0.57: 15,
0.58: 21,
0.59: 13,
0.6: 24,
0.61: 16,
0.62: 16,
0.63: 13,
0.64: 13,
0.65: 9,
0.66: 17,
0.67: 18,
0.68: 8,
0.69: 9,
0.7: 17,
0.71: 15,
0.72: 9,
0.73: 13,
0.74: 20,
0.75: 7,
0.76: 11,
0.77: 9,
0.78: 9,
0.79: 13,
0.8: 13,
0.81: 8,
0.82: 4,
0.83: 13,
0.84: 10,
0.85: 6,
0.86: 5,
0.87: 8,
0.88: 12,
0.89: 3,
0.9: 2,
0.91: 6,

0.92: 7,
0.93: 4,
0.94: 5,
0.95: 4,
0.96: 8,
0.97: 3,
0.98: 6,
0.99: 6,
1.0: 0.00023090169067000628,
1.01: 3,
1.02: 2,
1.03: 3,
1.04: 6,
1.05: 7,
1.06: 6,
1.07: 2,
1.08: 5,
1.09: 1,
1.11: 1,
1.12: 1,
1.13: 5,
1.14: 1,
1.15: 1,
1.16: 3,
1.17: 5,
1.18: 5,
1.19: 4,
1.2: 2,
1.21: 2,
1.22: 3,
1.23: 6,
1.24: 2,
1.26: 5,
1.27: 9,
1.28: 4,
1.29: 3,
1.3: 1,
1.31: 1,
1.32: 4,
1.33: 1,
1.34: 1,
1.36: 2,
1.37: 2,
1.38: 1,
1.39: 2,
1.4: 1,
1.45: 3,

1.46: 3,
1.48: 2,
1.49: 1,
1.5: 1,
1.51: 1,
1.52: 1,
1.53: 1,
1.54: 2,
1.55: 2,
1.56: 1,
1.6: 2,
1.61: 3,
1.63: 3,
1.64: 2,
1.65: 2,
1.66: 2,
1.67: 2,
1.68: 1,
1.7: 1,
1.75: 2,
1.76: 1,
1.78: 2,
1.83: 2,
1.85: 2,
1.93: 1,
2: 0.00010331872260488416,
2.04: 1,
2.08: 1,
2.14: 1,
2.18: 1,
2.23: 1,
2.26: 1,
2.58: 1,
2.7: 1,
2.72: 1,
2.98: 1,
3: 7.827175954915466e-05,
3.41: 1,
4: 0.00010644959298685035,
5: 0.00010488415779586725,
6: 8.140262993112086e-05,
7: 6.105197244834064e-05,
8: 6.966186599874766e-05,
9: 6.809643080776456e-05,
10: 5.165936130244208e-05,
11: 5.9486537257357546e-05,
12: 5.5572949279899807e-05,

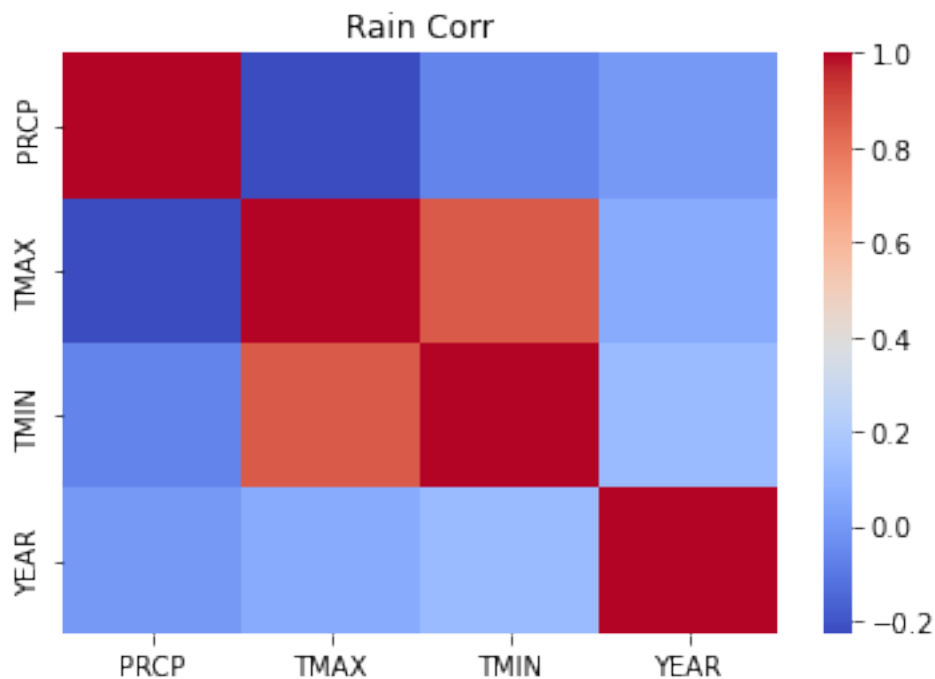
13: 5.8703819661866e-05,
14: 5.3224796493425174e-05,
15: 4.8528490920475895e-05,
16: 3.757044458359424e-05,
17: 4.53976205385097e-05,
18: 5.0876643706950534e-05,
19: 4.618033813400125e-05,
20: 4.304946775203507e-05,
22: 4.148403256105198e-05,
23: 4.226675015654352e-05,
25: 3.8353162179085784e-05,
27: 3.678772698810269e-05,
28: 3.209142141515341e-05,
29: 3.913587977457733e-05,
30: 3.6005009392611145e-05,
31: 3.287413901064496e-05,
32: 3.52222917971196e-05,
33: 3.443957420162805e-05,
35: 3.052598622417032e-05,
38: 2.739511584220413e-05,
40: 2.582968065122104e-05,
43: 2.8960551033187226e-05,
44: 2.4264245460237948e-05,
45: 2.9743268628678773e-05,
51: 2.3481527864746398e-05,
52: 2.6612398246712587e-05,
54: 2.113337507827176e-05,
56: 2.269881026925485e-05,
57: 1.8002504696305572e-05,
61: 2.5046963055729494e-05,
68: 2.0350657482780215e-05,
69: 1.643706950532248e-05,
71: 1.7219787100814026e-05,
72: 1.5654351909830933e-05,
88: 1.4871634314339386e-05,
90: 1.408891671884784e-05,
94: 1.3306199123356293e-05,
111: 1.1740763932373199e-05,
112: 1.0958046336881654e-05,
113: 1.0175328741390108e-05,
118: 1.2523481527864747e-05,
124: 9.39261114589856e-06,
132: 7.827175954915466e-06,
136: 7.04445835942392e-06,
157: 5.479023168440827e-06,
161: 6.2617407639323735e-06,
199: 4.69630557294928e-06,

```
207: 3.913587977457733e-06,
230: 3.1308703819661867e-06,
235: 2.34815278647464e-06,
354: 1.5654351909830934e-06,
471: 7.827175954915467e-07,
7403: 0.0}
```

```
[27]: #Create 1 CDF with one of your variables, using page 41-44 as your guide, what
      ↪ does this tell you about your variable and how does it address the question
      ↪ you are trying to answer (Chapter 4).
```

```
[29]: sns.heatmap(rain.corr(), cmap='coolwarm')
      plt.title('Rain Corr')
```

```
[29]: Text(0.5, 1.0, 'Rain Corr')
```



```
[30]: rain_per_year = rain[(rain['RAIN']==True)].groupby('YEAR').count()
```

```
[31]: year_rain_count=rain[rain['RAIN']==True].groupby('YEAR').count()
```

```
[32]: rain_per_year['RAIN']=year_rain_count['RAIN']
```

```
[33]: rain_per_year['TMIN MEAN']=rain.groupby('YEAR')['TMIN'].mean()
      rain_per_year['TMAX MEAN']=rain.groupby('YEAR')['TMAX'].mean()
      rain_per_year['PRCP MEAN']=rain.groupby('YEAR')['PRCP'].mean()
```

```
[ ]:
```

```
[34]: rain_per_year.reset_index(inplace=True)
```

```
[35]: rain_per_year = rain_per_year[['YEAR', 'RAIN', 'TMIN MEAN', 'TMAX MEAN', 'PRCP_
      ↪MEAN']]
```

```
[36]: rain_per_year.head()
```

```
[36]:
```


	YEAR	RAIN	TMIN MEAN	TMAX MEAN	PRCP MEAN
0	1948	177	41.196721	57.013661	0.125109
1	1949	139	41.391781	59.147945	0.088932
2	1950	194	41.000000	57.035616	0.151068
3	1951	150	41.052055	58.545205	0.110411
4	1952	139	41.467213	58.743169	0.064973

```
[37]: rain_per_year.columns
```

```
[37]: Index(['YEAR', 'RAIN', 'TMIN MEAN', 'TMAX MEAN', 'PRCP MEAN'], dtype='object')
```


```
[38]: sns.lmplot(x='YEAR', y='RAIN', data=rain_per_year)
```

```
[38]: <seaborn.axisgrid.FacetGrid at 0x2b720ec2a90>
```




```
[39]: sns.jointplot(x='YEAR', y='RAIN', data=rain_per_year, kind='kde')
```

```
[39]: <seaborn.axisgrid.JointGrid at 0x2b720e14d00>
```




```
[40]: plt.plot('YEAR', 'TMIN MEAN', data=rain_per_year, marker='o')  
      plt.plot('YEAR', 'TMAX MEAN', data=rain_per_year, marker='', color='olive')
```

```
[40]: [<matplotlib.lines.Line2D at 0x2b721153160>]
```

```
[41]: plt.plot('YEAR', 'TMIN', data=rain, marker='o')
      plt.plot('YEAR', 'TMAX', data=rain, marker='', color='olive')
```

```
[41]: [<matplotlib.lines.Line2D at 0x2b7211afb80>]
```




```
[42]: sns.lmplot('YEAR', 'PRCP MEAN', data=rain_per_year)
```

```
C:\Users\mgan\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variables as keyword args: x, y. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
warnings.warn(  

```

```
[42]: <seaborn.axisgrid.FacetGrid at 0x2b7211dd1f0>
```




```
[43]: sns.lmplot('YEAR', 'PRCP', data=rain)
```

```
C:\Users\mgan\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variables as keyword args: x, y. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
warnings.warn(  

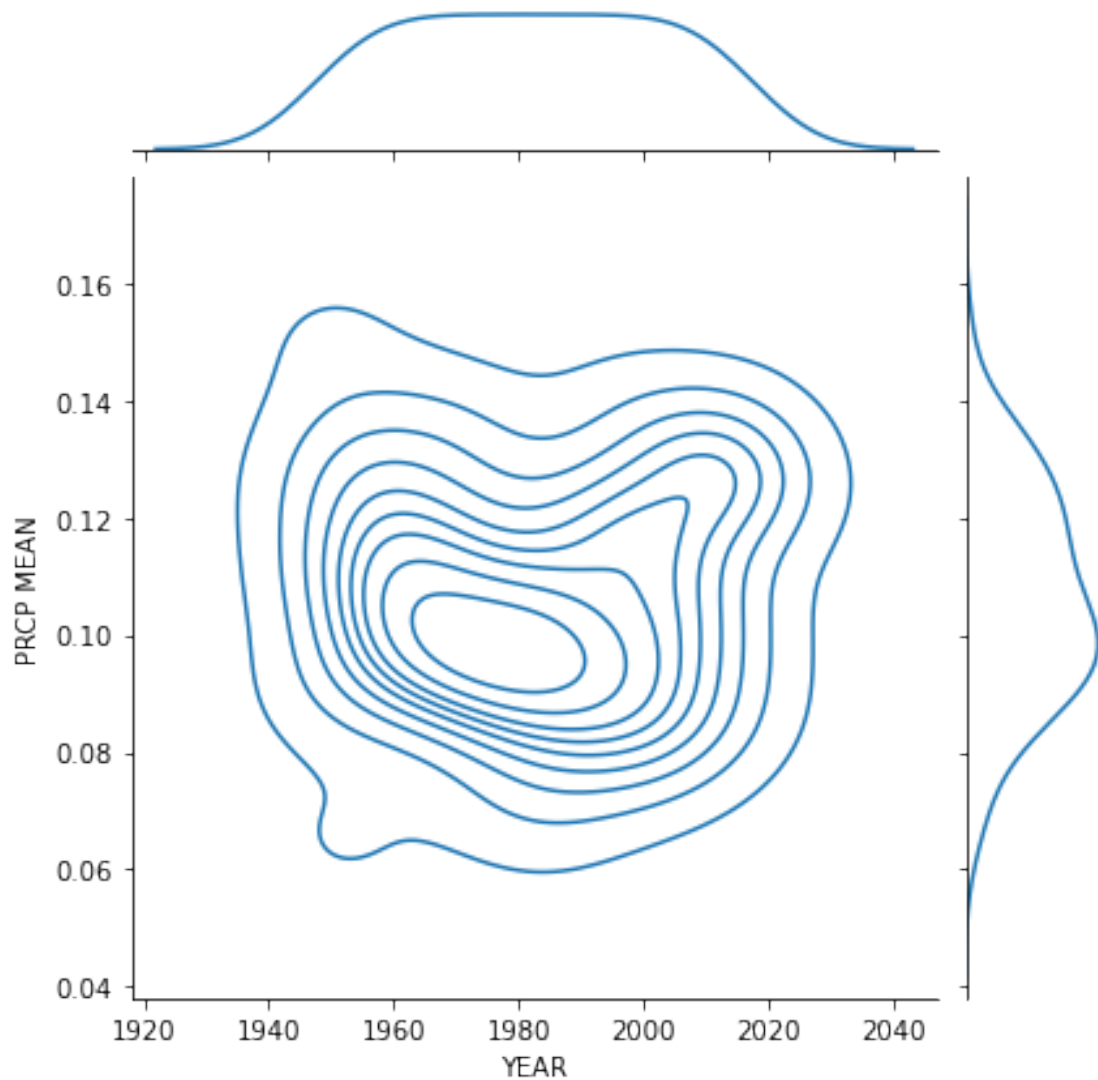
```

```
[43]: <seaborn.axisgrid.FacetGrid at 0x2b7211d6b50>
```



```
[44]: sns.jointplot(x='YEAR', y='PRCP MEAN', data=rain_per_year, kind='kde')
```

```
[44]: <seaborn.axisgrid.JointGrid at 0x2b721232520>
```



```
[50]: super_column = rain_per_year[['YEAR', 'TMIN MEAN', 'TMAX MEAN']]
```

```
[51]: super_column['Max'] = 'Max'
```

C:\Users\mgghan\AppData\Local\Temp\ipykernel_1592\3469924724.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
super_column['Max'] = 'Max'
```

```
[52]: super_column['Min'] = 'Min'
```

```
[53]: max = super_column[['YEAR', 'TMAX MEAN', 'Max']]
      min = super_column[['YEAR', 'TMIN MEAN', 'Min']]
```

```
[54]: max['temp']=max['TMAX MEAN']
      max['type']=max['Max']
      min['temp']=min['TMIN MEAN']
      min['type']=min['Min']
```

C:\Users\mghan\AppData\Local\Temp\ipykernel_1592\3706309177.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
max['temp']=max['TMAX MEAN']
```

C:\Users\mghan\AppData\Local\Temp\ipykernel_1592\3706309177.py:3:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
min['temp']=min['TMIN MEAN']
```

```
[55]: max = max[['YEAR', 'temp', 'type']]
      min = min[['YEAR', 'temp', 'type']]
```

```
[56]: combined = max.append(min)
```

C:\Users\mghan\AppData\Local\Temp\ipykernel_1592\299298275.py:1: FutureWarning:


The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
combined = max.append(min)
```

```
[ ]:
```


```
[57]: sns.heatmap(rain_per_year.corr(), cmap='coolwarm')
```

```
[57]: <AxesSubplot:>
```




```
[58]: sns.lmplot(x='YEAR', y='temp', hue='type', data=combined)
```


```
[58]: <seaborn.axisgrid.FacetGrid at 0x2b7224527f0>
```



```
[59]: sns.lmplot(x='YEAR', y='TMIN MEAN', data=rain_per_year)
      sns.lmplot(x='YEAR', y='TMAX MEAN', data=rain_per_year)
```

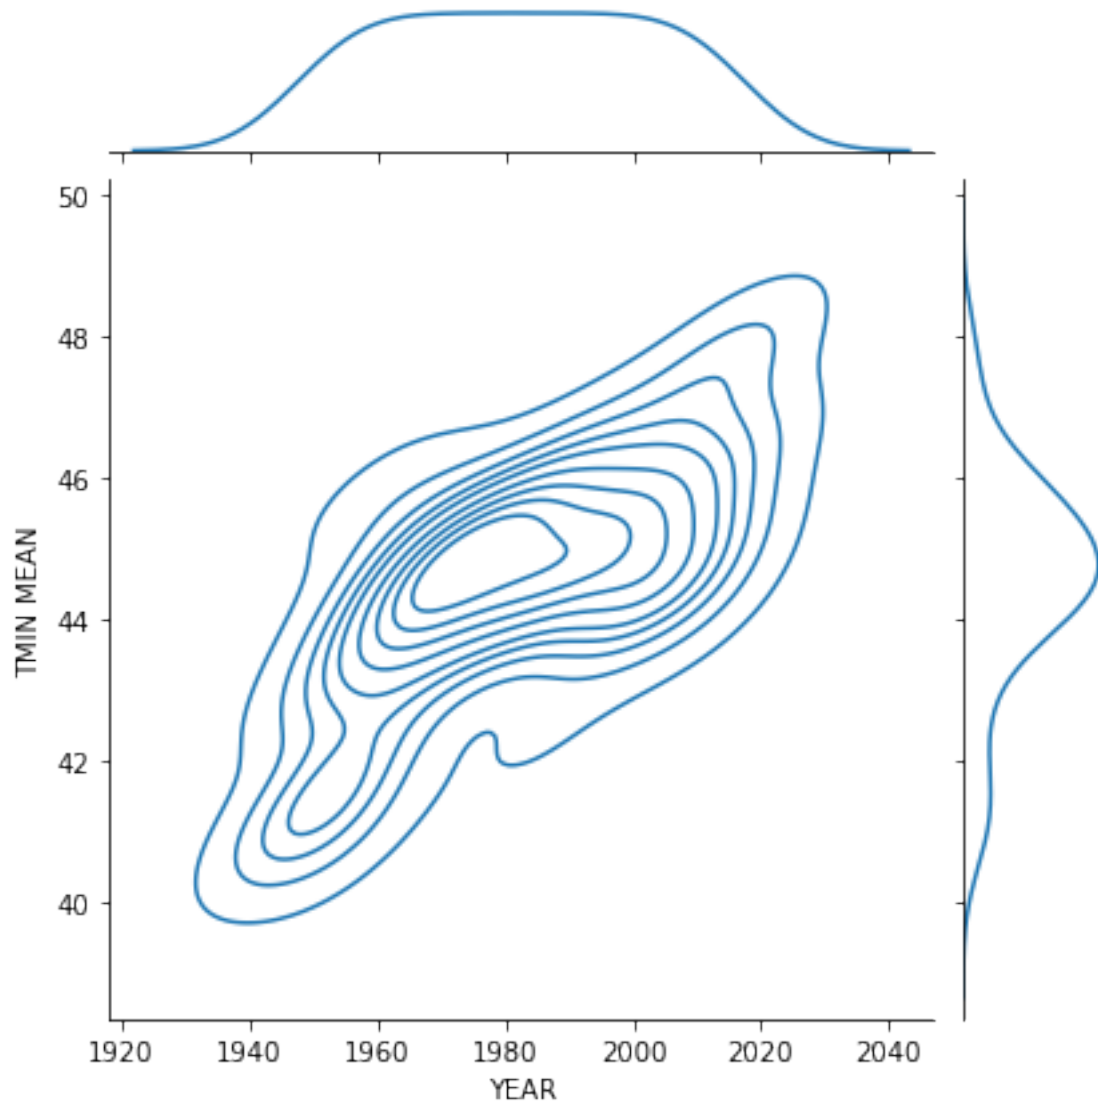
```
[59]: <seaborn.axisgrid.FacetGrid at 0x2b7223c47c0>
```






```
[60]: sns.jointplot(x='YEAR', y='TMIN MEAN', data=rain_per_year, kind='kde')  
sns.jointplot(x='YEAR', y='TMAX MEAN', data=rain_per_year, kind='kde')
```


```
[60]: <seaborn.axisgrid.JointGrid at 0x2b7226bc580>
```





```
[61]: sns.lmplot(x='YEAR', y='RAIN', data=rain_per_year)
```

```
[61]: <seaborn.axisgrid.FacetGrid at 0x2b722957340>
```



```
[62]: import statsmodels.formula.api as sm
```

```
[63]: result = sm.ols(formula='RAIN ~ YEAR',data=rain_per_year).fit()
print(result.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          RAIN    R-squared:                0.022
Model:                  OLS     Adj. R-squared:           0.007
Method:                 Least Squares    F-statistic:            1.520
Date:                   Fri, 31 May 2024    Prob (F-statistic):      0.222
Time:                   18:34:13    Log-Likelihood:         -289.74
No. Observations:       70        AIC:                   583.5
Df Residuals:           68        BIC:                   588.0
Df Model:                1
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
YEAR	-0.001	0.001	-0.5	0.62	-0.002	0.001
Intercept	159.5	3.5	45.5	<0.001	152.5	166.5

```

-----
Intercept    378.4702    180.659    2.095    0.040    17.970    738.970
YEAR         -0.1124     0.091    -1.233    0.222    -0.294    0.069
=====
Omnibus:                0.144    Durbin-Watson:                1.891
Prob(Omnibus):          0.931    Jarque-Bera (JB):            0.196
Skew:                   0.102    Prob(JB):                    0.907
Kurtosis:               2.839    Cond. No.                    1.95e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
[64]: result = sm.ols(formula='RAIN ~ YEAR',data=rain_per_year).fit()
print(result.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          RAIN    R-squared:                0.022
Model:                  OLS    Adj. R-squared:         0.007
Method:                 Least Squares    F-statistic:          1.520
Date:                  Fri, 31 May 2024    Prob (F-statistic):    0.222
Time:                  18:34:17    Log-Likelihood:       -289.74
No. Observations:      70    AIC:                  583.5
Df Residuals:          68    BIC:                  588.0
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	378.4702	180.659	2.095	0.040	17.970	738.970
YEAR	-0.1124	0.091	-1.233	0.222	-0.294	0.069

```

=====
Omnibus:                0.144    Durbin-Watson:                1.891
Prob(Omnibus):          0.931    Jarque-Bera (JB):            0.196
Skew:                   0.102    Prob(JB):                    0.907
Kurtosis:               2.839    Cond. No.                    1.95e+05
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.95e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
[65]: result.pvalues
```

```
[65]: Intercept    0.039904  
      YEAR        0.221791  
      dtype: float64
```

```
[ ]:
```